

COMP9311 (Database Systems)

Week 5

PLpgSQL

Procedural Language Extensions for PostgreSQL

Limitations of Basic SQL

- What we have seen of SQL so far:
 - data definition language (*create table(...)*)
 - constraints (*domain, key, referential integrity*)
 - query language (*select...from...where...*)
 - views (give names to SQL queries)
- This is not sufficient to write complete applications.
- More **extensibility** and **programmability** are needed.

Extending SQL

- Ways in which standard SQL might be extended:
 - new data types (incl. constraints, I/O, indexes, ...)
 - object-orientation
 - more powerful constraint checking
 - packaging/parameterizing queries
 - more functions/aggregates for use in queries
 - event-based triggered actions
 - massive data, spread over a network
- All are required to assist in application development.

SQL Data Types

- SQL data definition language provides:
 - atomic types: integer, float, character, boolean
 - ability to define tuple types (*create table*)
- SQL also provides mechanisms to define new types:
 - basic types: CREATE DOMAIN
 - tuple types: CREATE TYPE

SQL Data Types_(cont.)

- Defining an atomic type (as specialisation of existing type):

```
CREATE DOMAIN DomainName [ AS ] DataType  
[ DEFAULT expression ]  
[ CONSTRAINT ConstrName constraint ]
```

- Example

```
create domain UnswCourseCode as text  
check ( value ~ '[A - Z ]{4}[0 -9]{4} ' );
```

- which can then be used like other SQL atomic types, e.g.

```
create table Course (  
    id integer ,  
    code UnswCourseCode ,  
    ...  
);
```

SQL Data Types_(cont.)

- Defining a tuple type:

```
CREATE TYPE TypeName AS  
( AttrName1 DataType1 , AttrName2 DataType2 , ...)
```

- Example

```
create type ComplexNumber as ( r float , i float );  
  
create type CourseInfo as (  
    course UnswCourseCode ,  
    syllabus text ,  
    lecturer text  
);
```

- If attributes need constraints, can be supplied by using a DOMAIN.

SQL Data Types_(cont.)

- Other ways that tuple types are defined in SQL:
 - CREATE TABLE T (effectively creates tuple type T)
 - CREATE VIEW V (effectively creates tuple type V)
- CREATE TYPE is different from CREATE TABLE:
 - does not create a new (empty) table
 - does not provide for key constraints
 - does not have explicit specification of domain constraints
- Used for specifying return types of functions that return tuples or sets.

SQL as a Programming Language

- SQL is a powerful language for manipulating relational data. But it is not a powerful programming language.
- At some point in developing complete database applications
 - we need to implement user interactions
 - we need to control sequences of database operations
 - we need to process query results in complex ways
- and SQL cannot do any of these.
- SQL cannot even do something as simple as factorial

What's wrong with SQL?

- Consider the problem of withdrawal from a bank account:
- If a bank customer attempts to withdraw more funds than they have in their account, then indicate 'Insufficient Funds', otherwise update the account.
- An attempt to implement this in SQL

What's wrong with SQL? (cont.)

- Solution:

```
select ' Insufficient Funds '  
from Accounts  
where acctNo = AcctNum and balance < Amount;
```

```
update Accounts  
set balance = balance - Amount  
where acctNo = AcctNum and balance >= Amount;
```

```
select ' New balance : ' || balance  
from Accounts  
where acctNo = AcctNum;
```

What's wrong with SQL? (cont.)

- Two possible evaluation scenarios:
 - displays 'Insufficient Funds', UPDATE has no effect, displays unchanged balance
 - UPDATE occurs as required, displays changed balance

What's wrong with SQL? (cont.)

- Some problems:
 - SQL doesn't allow parameterisation (e.g. *AcctNum*)
 - always attempts UPDATE, even when it knows it's invalid
 - always displays balance, even when not changed
- To accurately express the “business logic”, we need facilities like conditional execution and parameter passing.

Database Programming_(cont.)

- Database programming requires a combination of
 - manipulation of data in DB (via SQL)
 - conventional programming (via procedural code)
- This combination is realised in a number of ways:
 - passing SQL commands via a "call-level" interface
(PL is decoupled from DBMS; most flexible; e.g. Java/JDBC, PHP)
 - embedding SQL into augmented programming languages
(requires PL pre-processor; typically DBMS-specific; e.g. SQL/C)
 - special-purpose programming languages in the DBMS
(integrated with DBMS; enables extensibility; e.g. PL/SQL, PLpgSQL)

Database Programming_(cont.)

- **Recap the example:**
- withdraw *amount* dollars from account *acctNum*
- using a function with parameters *amount* and *acctNum*
- returning two possible text results :
 - 'Insufficient funds' if try to withdraw too much
 - 'New balance *newAmount*' if withdrawal ok
- an obvious side-effect is to change the stored balance
- Requires a combination of
 - SQL code to access the database
 - procedural code to control the process

Database Programming_(cont.)

Stored-procedure approach (PLpgSQL):

```
create function
    withdraw(acctNum text, amount integer) returns text as $$
declare bal integer;
begin
    select balance into bal
    from Accounts
    where acctNo = acctNum;
    if (bal < amount) then
        return 'Insufficient Funds';
    else
        update Accounts
        set balance = balance - amount
        where acctNo = acctNum;
        select balance into bal
        from Accounts where acctNo = acctNum;
        return 'New Balance: ' || bal;
    end if;
end;
$$ language plpgsql;
```

Stored Procedures

- Stored procedures
 - procedures/functions that are stored in DB along with data
 - written in a language combining SQL and procedural ideas
 - provide a way to extend operations available in database
 - executed within the DBMS (close coupling with query engine)
- Benefits of using stored procedures:
 - minimal data transfer cost SQL \leftrightarrow procedural code
 - user-defined functions can be nicely integrated with SQL
 - procedures are managed like other DBMS data (ACID)
 - procedures and the data they manipulate are held together

SQL Functions

- PostgreSQL Manual: 37.4. Query Language (SQL) Functions (see 37.4 if using version 10.5)
- PostgreSQL allows functions to be defined in SQL

CREATE OR REPLACE FUNCTION

funcName(arg1type, arg2type,)

RETURNS *rettype*

AS \$\$

SQL statements

\$\$ LANGUAGE sql;

SQL Functions_(cont.)

- Within the function, arguments are accessed as \$1, \$2, ...
- Return value: result of the last SQL statement.
- *rettype* can be any PostgreSQL data type (incl tuples, tables).
- Function returning a table: returns set of *TupleType*

SQL Functions_(cont.)

- Examples:

-- max price of specified beer

create or replace function

maxPrice(text) returns float

as \$\$

select max(price) from Sells where beer = \$1;

\$\$ language sql;

SQL Functions_(cont.)

-- usage examples

```
select maxPrice('New');
```

```
maxprice
```

```
-----
```

```
2.8
```

```
select bar,price from sells
```

```
where beer='New' and price=maxPrice('New');
```

```
bar
```

```
price
```

```
-----
```

```
-----
```

```
Marble Bar
```

```
2.8
```

SQL Functions_(cont.)

- Examples:

-- set of Bars from specified suburb

create or replace function

 hotelsIn(text) returns setof Bars

as \$\$

 select * from Bars where addr = \$1;

\$\$ language sql;

SQL Functions_(cont.)

-- usage examples

```
select * from hotelsIn('The Rocks');
```

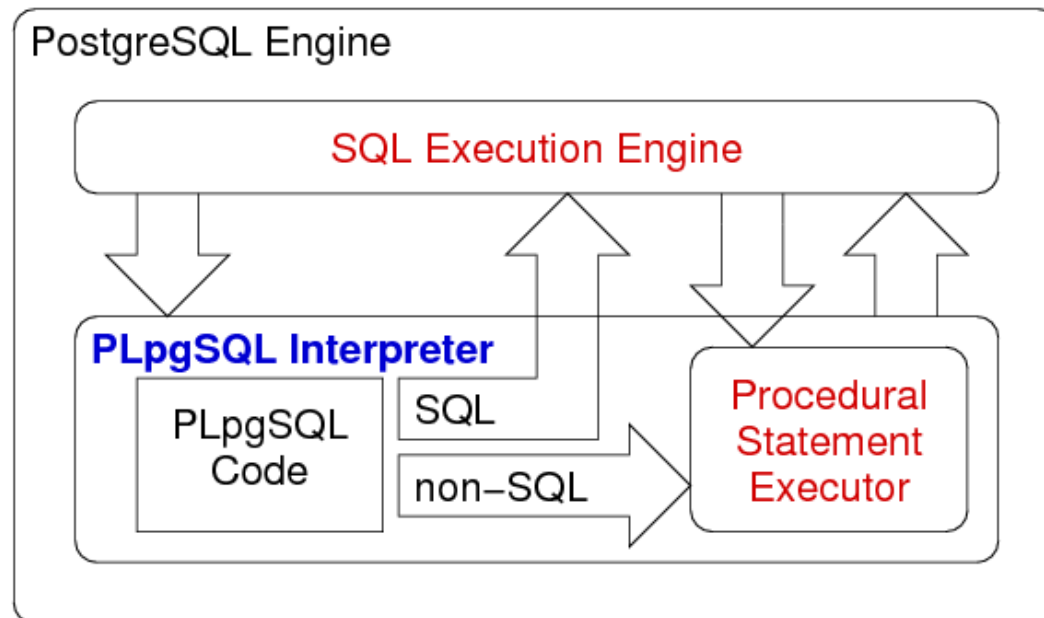
name	addr	license
-----	-----	-----
Australia Hotel	The Rocks	123456
Lord Nelson	The Rocks	123888

PLpgSQL

- PostgreSQL Manual: Chapter 42: PLpgSQL (if using version 10.5)
- PLpgSQL = **P**rocedural **L**anguage extensions to **P**ostgre**SQL**
- A PostgreSQL-specific language integrating features of:
 - procedural programming and SQL programming
- Functions are stored in the database with the data.
- Provides a means for extending DBMS functionality, e.g.
 - implementing constraint checking (triggered functions)
 - complex query evaluation (e.g. recursive)
 - complex computation of column values
 - detailed control of displayed results

PLpgSQL_(cont)

- The PLpgSQL interpreter
 - executes procedural code and manages variables
 - calls PostgreSQL engine to evaluate SQL statements



Defining PLpgSQL Functions

- PLpgSQL functions are created (and inserted into db) via:

```
CREATE OR REPLACE
    funcName(param1, param2, ....)
    RETURNS rettype
AS $$
DECLARE
    variable declarations
BEGIN
    code for function
END;
$$ LANGUAGE plpgsql;
```

- Note: the entire function body is a single SQL string.

Defining PLpgSQL Functions_(cont.)

Recap Stored-procedure approach (PLpgSQL):

```
create function
    withdraw(acctNum text, amount integer) returns text as $$
declare bal integer;
begin
    select balance into bal
    from Accounts
    where acctNo = acctNum;
    if (bal < amount) then
        return 'Insufficient Funds';
    else
        update Accounts
        set balance = balance - amount
        where acctNo = acctNum;
        select balance into bal
        from Accounts where acctNo = acctNum;
        return 'New Balance: ' || bal;
    end if;
end;
$$ language plpgsql;
```

PLpgSQL Function Parameters

- All parameters are passed by value in PLpgSQL.
- Within a function, parameters can be referred to:
 - using positional notation (\$1, \$2, ...)
 - via aliases, supplied either
 - as part of the function header (e.g. f(a int, b int))
 - as part of the declarations (e.g. a alias for \$1; b alias for \$2)

PLpgSQL Function Parameters_(cont.)

- Example: old-style function

```
CREATE OR REPLACE FUNCTION
    cat(text, text) RETURNS text
AS '
DECLARE
    x alias for $1; -- alias for parameter
    y alias for $2; -- alias for parameter
    result text; -- local variable
BEGIN
    result := x || '||' || y;
    return result;
END;
' LANGUAGE 'plpgsql';
```

- **Beware:** never give aliases the same names as attributes.

PLpgSQL Function Parameters_(cont.)

- Example: new-style function

```
CREATE OR REPLACE FUNCTION
    add(x text, y text) RETURNS text
AS $$
DECLARE
    result text; -- local variable
BEGIN
    result := x || '|' || y;
    return result;
END;
$$ LANGUAGE 'plpgsql';
```

- **Beware:** never give aliases the same names as attributes.

PLpgSQL Function Parameters_(cont.)

```
CREATE OR REPLACE FUNCTION
    add ( x anyelement , y anyelement ) RETURNS anyelement
AS $$
BEGIN
    return x + y ;
END ;
$$ LANGUAGE plpgsql ;
```

- Restrictions: requires x and y to have values of the same "addable" type

PLpgSQL Function Parameters_(cont.)

- PLpgSQL allows overloading (i.e. same name, different arg types)
- Example

```
CREATE FUNCTION add ( int , int ) RETURNS int AS  
$$ BEGIN return $1 + $2 ; END ; $$ LANGUAGE plpgsql ;
```

```
CREATE FUNCTION add ( int , int , int ) RETURNS int AS  
$$ BEGIN return $1 + $2 + $3 ; END ; $$ LANGUAGE plpgsql ;
```

```
CREATE FUNCTION add ( char (1) , int ) RETURNS int AS  
$$ BEGIN return ascii ( $1 )+ $2 ; END ; $$ LANGUAGE plpgsql ;
```

- But must differ in arg types, so cannot also define:

```
CREATE FUNCTION add ( char (1) , int ) RETURNS char AS  
$$ BEGIN return chr ( ascii ( $1 )+ $2 ); END ; $$ LANGUAGE plpgsql ;
```

- i.e. cannot have two functions that look like add(char(1), int).

Function Return Types

- A PostgreSQL function can return a value which is
 - void (i.e. no return value)
 - an atomic data type (e.g. integer, text, ...)
 - a tuple (e.g. table record type or tuple type)
 - a set of atomic values (like a table column)
 - a set of tuples (i.e. a table)
- A function returning a set of tuples is similar to a view.

Function Return Types_(cont)

- Examples of different function return types:

```
create type Employee as  
  (id integer, name text, salary float, ...);
```

```
create function factorial(integer)  
  returns integer ...
```

```
create function EmployeeOfMonth(date)  
  returns Employee ...
```

```
create function allSalaries()  
  returns setof float ...
```

```
create function OlderEmployees()  
  returns setof Employee ...
```

Function Return Types_(cont)

- Different kinds of functions are invoked in different ways:

```
select factorial(5);
```

```
-- returns one integer
```

```
select EmployeeOfMonth('2008-04-01');
```

```
-- returns (x,y,z,...)
```

```
select * from EmployeeOfMonth('2008-04-01');
```

```
-- one-row table
```

```
select * from allSalaries();
```

```
-- single-column table
```

```
select * from OlderEmployees();
```

```
-- subset of Employees
```

Using PLpgSQL Functions

- PLpgSQL functions can be invoked in several ways:
 - as part of a SELECT statement

```
select myFunction ( arg1 , arg2 );
```

```
select * from myTableFunction ( arg1 , arg2 );
```
 - as part of the execution of another PLpgSQL function

```
PERFORM myVoidFunction ( arg1 , arg2 );
```

```
result := myOtherFunction ( arg1 );
```
 - automatically, via an insert/delete/update trigger

```
create trigger T before update on R
```

```
for each row execute procedure myCheck ();
```

Special Data Types

- by deriving a type from an existing database table, e.g.

account Accounts % ROWTYPE ;

- Record components referenced via attribute name account.branchName%TYPE

Special Data Types_(cont.)

- Variables can also be defined in terms of:
 - the type of an existing variable or table column
 - the type of an existing table row (implicit RECORD type)

- Example

quantity INTEGER ;

start_qty quantity % TYPE ;

employee Employees % ROWTYPE ;

name Employees.name % TYPE ;

Control Structures

- Assignment
 - variable := expression;
- Example:
 - tax := subtotal * 0.06;
 - my_record.user_id := 20;
- Conditionals
 - IF ... THEN
 - IF ... THEN ... ELSE
 - IF ... THEN ... ELSIF ... THEN ... ELSE
- Example
 - IF v_user_id > 0 THEN
 - UPDATE users SET email = v_email WHERE user_id = v_user_id; END IF;

Control Structures_(cont.)

- Iteration

```
LOOP
```

```
    Statement
```

```
END LOOP ;
```

- Example

```
LOOP
```

```
    IF count > 0 THEN
```

```
        -- some computations
```

```
    END IF;
```

```
END LOOP;
```

Control Structures_(cont.)

- Iteration

```
FOR int_var IN low .. high LOOP
```

```
    Statement
```

```
END LOOP ;
```

- Example

```
FOR i IN 1..10 LOOP
```

```
    -- i will take on the values 1,2,3,4,5,6,7,8,9,10 within the  
    loop
```

```
END LOOP;
```


SELECT ... INTO

- Can capture query results via:
SELECT Exp1 , Exp2 , ... , Expn
INTO Var1 , Var2 , ... , Varn
FROM TableList
WHERE Condition ...
- The semantics:
- execute the query as usual
- return “projection list” (Exp1, Exp2, ...) as usual
- assign each Exp_i to corresponding Vari

SELECT ... INTO_(cont.)

- Assigning a simple value via SELECT ... INTO:

-- cost is local var , price is attr

SELECT price INTO cost

FROM StockList

WHERE item = ' Cricket Bat ';

cost := cost * (1 + tax_rate);

total := total + cost ;

Exceptions

- Syntax

BEGIN

Statements ...

EXCEPTION

WHEN Exceptions1 THEN

StatementsForHandler1

WHEN Exceptions2 THEN

StatementsForHandler2

...

END ;

- Each Exceptions_i is an OR list of exception names, e.g.,
 - `division_by_zero` OR `floating_point_exception` OR ...

Exceptions_(cont.)

- Example

-- table T contains one tuple (' Tom ' , ' Jones ')

DECLARE

 x INTEGER := 3;

BEGIN

 UPDATE T SET firstname = ' Joe ' WHERE lastname = ' Jones ';

 -- table T now contains (' Joe ' , ' Jones ')

 x := x + 1;

 y := x / y; ---- y: = # of Tom Jones in Staff Table

EXCEPTION

 WHEN division_by_zero THEN

 -- update on T is rolled back to (' Tom ' , ' Jones ')

 RAISE NOTICE ' Caught division_by_zero ';

 RETURN x ;

 -- value returned is 4

END ;

Exceptions_(cont.)

- The RAISE operator generates server log entries, e.g.
 - RAISE DEBUG ' Simple message ';
 - RAISE NOTICE ' User = % ' , user_id ;
 - RAISE EXCEPTION ' Fatal : value was % ' , value ;
- There are several levels of severity:
 - DEBUG, LOG, INFO, NOTICE, WARNING, and EXCEPTION
 - not all severities generate a message to the client

Cursors

- A cursor is a variable that can be used to access the result of a particular SQL query
- Cursors move sequentially from row to row (cf., file pointers in C).

Employees

Id	Name	Salary
961234	John Smith	35000.00
954321	Kevin Smith	48000.00
912222	David Smith	31000.00

cursor - - ->

Cursors_(cont.)

- Simplest way to use cursors: implicitly via FOR ... IN
- Requires: RECORD variable or Table%ROWTYPE variable
- Example:

```
CREATE FUNCTION totalsal () RETURNS REAL AS $$  
DECLARE  
    emp RECORD ;  
    total REAL := 0;  
BEGIN  
    FOR emp IN SELECT * FROM Employees  
    LOOP  
        total := total + emp . salary ;  
    END LOOP ;  
    RETURN total ;  
END ; $$ LANGUAGE plpgsql ;
```

- This style accounts for 95% of cursor usage.

Cursors_(cont.)

- Of course, the previous example would be better done as:

```
CREATE FUNCTION totalsal () RETURNS REAL AS $$
```

```
DECLARE
```

```
    total REAL ;
```

```
BEGIN
```

```
    SELECT sum ( salary ) INTO total FROM Employees ;
```

```
    return total ;
```

```
END ; $$ LANGUAGE plpgsql ;
```

- The iteration/summation can be done much more efficiently as an aggregation.

Cursors_(cont.)

- Basic operations on cursors: OPEN, FETCH, CLOSE

```
-- assume ... e CURSOR FOR SELECT * FROM Employees ;
```

```
OPEN e ;
```

```
LOOP
```

```
    FETCH e INTO emp ;
```

```
    EXIT WHEN NOT FOUND ;
```

```
    total := total + emp.salary ;
```

```
END LOOP ;
```

```
CLOSE e ;
```

Cursors_(cont.)

- The FETCH operation can also extract components of a row:

```
FETCH e INTO my_id , my_name , my_salary ;
```

- There must be one variable, of the correct type, for each column in the result.

Triggers

- Triggers are
 - procedures stored in the database
 - activated in response to database events (e.g., updates)
- Examples of uses for triggers:
 - maintaining summary data
 - checking schema-level constraints (assertions) on update
 - performing multi-table updates (to maintain assertions)

Triggers_(cont.)

- Triggers provide event-condition-action (ECA) programming:
 - an event activates the trigger
 - on activation, the trigger checks a condition
 - if the condition holds, a procedure is executed (the action)

Triggers_(cont.)

- Consider two triggers and an INSERT statement
create trigger X before insert on T Code1;
create trigger Y after insert on T Code2;
insert into T values (a,b,c,...);
- Consider two triggers and an UPDATE statement
create trigger X before update on T Code1;
create trigger Y after update on T Code2;
update T set b=j,c=k where a=m;

Triggers in PostgreSQL

- PostgreSQL triggers provide a mechanism for INSERT, DELETE or UPDATE events to automatically activate PLpgSQL functions
- Syntax for PostgreSQL trigger definition:

```
CREATE TRIGGER TriggerName  
{AFTER|BEFORE} Event1 [OR Event2 ...]  
ON TableName  
[ WHEN ( Condition ) ]  
FOR EACH {ROW|STATEMENT}  
EXECUTE PROCEDURE FunctionName(args...);
```

Triggers in PostgreSQL_(cont.)

- PLpgSQL Functions for Triggers

CREATE OR REPLACE FUNCTION name () RETURNS TRIGGER ..

- There is no restriction on what code can go in the function.
- However
 - RETURN OLD or RETURN new (depending on which version of the tuple is to be used)
 - Raise an EXCEPTION. In that case, no change occurs

Trigger Example

- Consider a database of people in the USA:
create table Person (
 id integer primary key,
 ssn varchar(11) unique,
 ... e.g. family, given, street, town ...
 state char(2), ...
);
create table States (
 id integer primary key,
 code char(2) unique,
 ... e.g. name, area, population, flag ...
);
- Constraint: $\text{Person.state} \in (\text{select code from States})$, or
exists (select id from States where code=Person.state)

Trigger Example_(cont.)

- **Example:** ensure that only valid state codes are used:

```
create trigger checkState before insert or update on Person for each row execute procedure  
checkState();
```

```
create function checkState() returns trigger as $$  
begin  
    -- normalise the user-supplied value  
    new.state = upper(trim(new.state));  
    if (new.state !~ '^[A-Z][A-Z]$') then  
        raise exception 'Code must be two alpha chars';  
    end if;  
    -- implement referential integrity check  
    select * from States where code=new.state;  
    if (not found) then  
        raise exception 'Invalid code %',new.state;  
    end if;  
    return new;  
end;  
$$ language plpgsql;
```

Trigger Example_(cont.)

- **Example:** department salary totals

- Scenario:

Employee(id, name, address, dept, salary, ...)

Department(id, name, manager, totSal, ...)

- An assertion that we wish to maintain:

Department.totSal =

(select sum(e.salary) from Employee e where
e.dept = d.id)))

Trigger Example_(cont.)

- Events that might affect the validity of the database
 - a new employee starts work in some department
 - an employee gets a rise in salary
 - an employee changes from one department to another
 - an employee leaves the company
- A single assertion could check validity after each change.
- With triggers, we have to program each case separately.
- Each program implements updates to *ensure* assertion holds.

Trigger Example_(cont.)

- Implement the Employee update triggers from above in PostgreSQL:
- Case 1: new employees arrive

```
create trigger TotalSalary1
after insert on Employees
for each row execute procedure totalSalary1();

create function totalSalary1() returns trigger
as $$
begin
    if (new.dept is not null) then
        update Department
        set totSal = totSal + new.salary
        where Department.id = new.dept;
    end if;
    return new;
end; $$ language plpgsql;
```

Trigger Example_(cont.)

- Case 2: employees change departments/salaries

```
create trigger TotalSalary2
```

```
after update on Employee
```

```
for each row execute procedure totalSalary2();
```

```
create function totalSalary2() returns trigger
```

```
as $$
```

```
begin
```

```
    update Department
```

```
    set totSal = totSal + new.salary
```

```
    where Department.id = new.dept;
```

```
    update Department set totSal = totSal - old.salary
```

```
    where Department.id = old.dept;
```

```
    return new;
```

```
end; $$ language plpgsql;
```

Trigger Example_(cont.)

- Case 3: employees leave

```
create trigger TotalSalary3  
after delete on Employee  
for each row execute procedure totalSalary3();
```

```
create function totalSalary3() returns trigger  
as $$  
begin  
    if (old.dept is not null) then  
        update Department  
        set totSal = totSal - old.salary where Department.id = old.dept;  
    end if;  
    return old;  
end; $$ language plpgsql;
```